

MapReduce

Lec -03

MapReduce

- MapReduce is a software framework and programming model used for processing huge amounts of data parallel by dividing the problem into some smaller and independent tasks.
- The **MapReduce** program works in two phases, namely
 - Map
 - Reduce.
- **Map tasks** deal with the splitting and mapping of data
- **Reduce tasks** shuffle and reduce the data. (Aggregate, summarize, filter or transform)

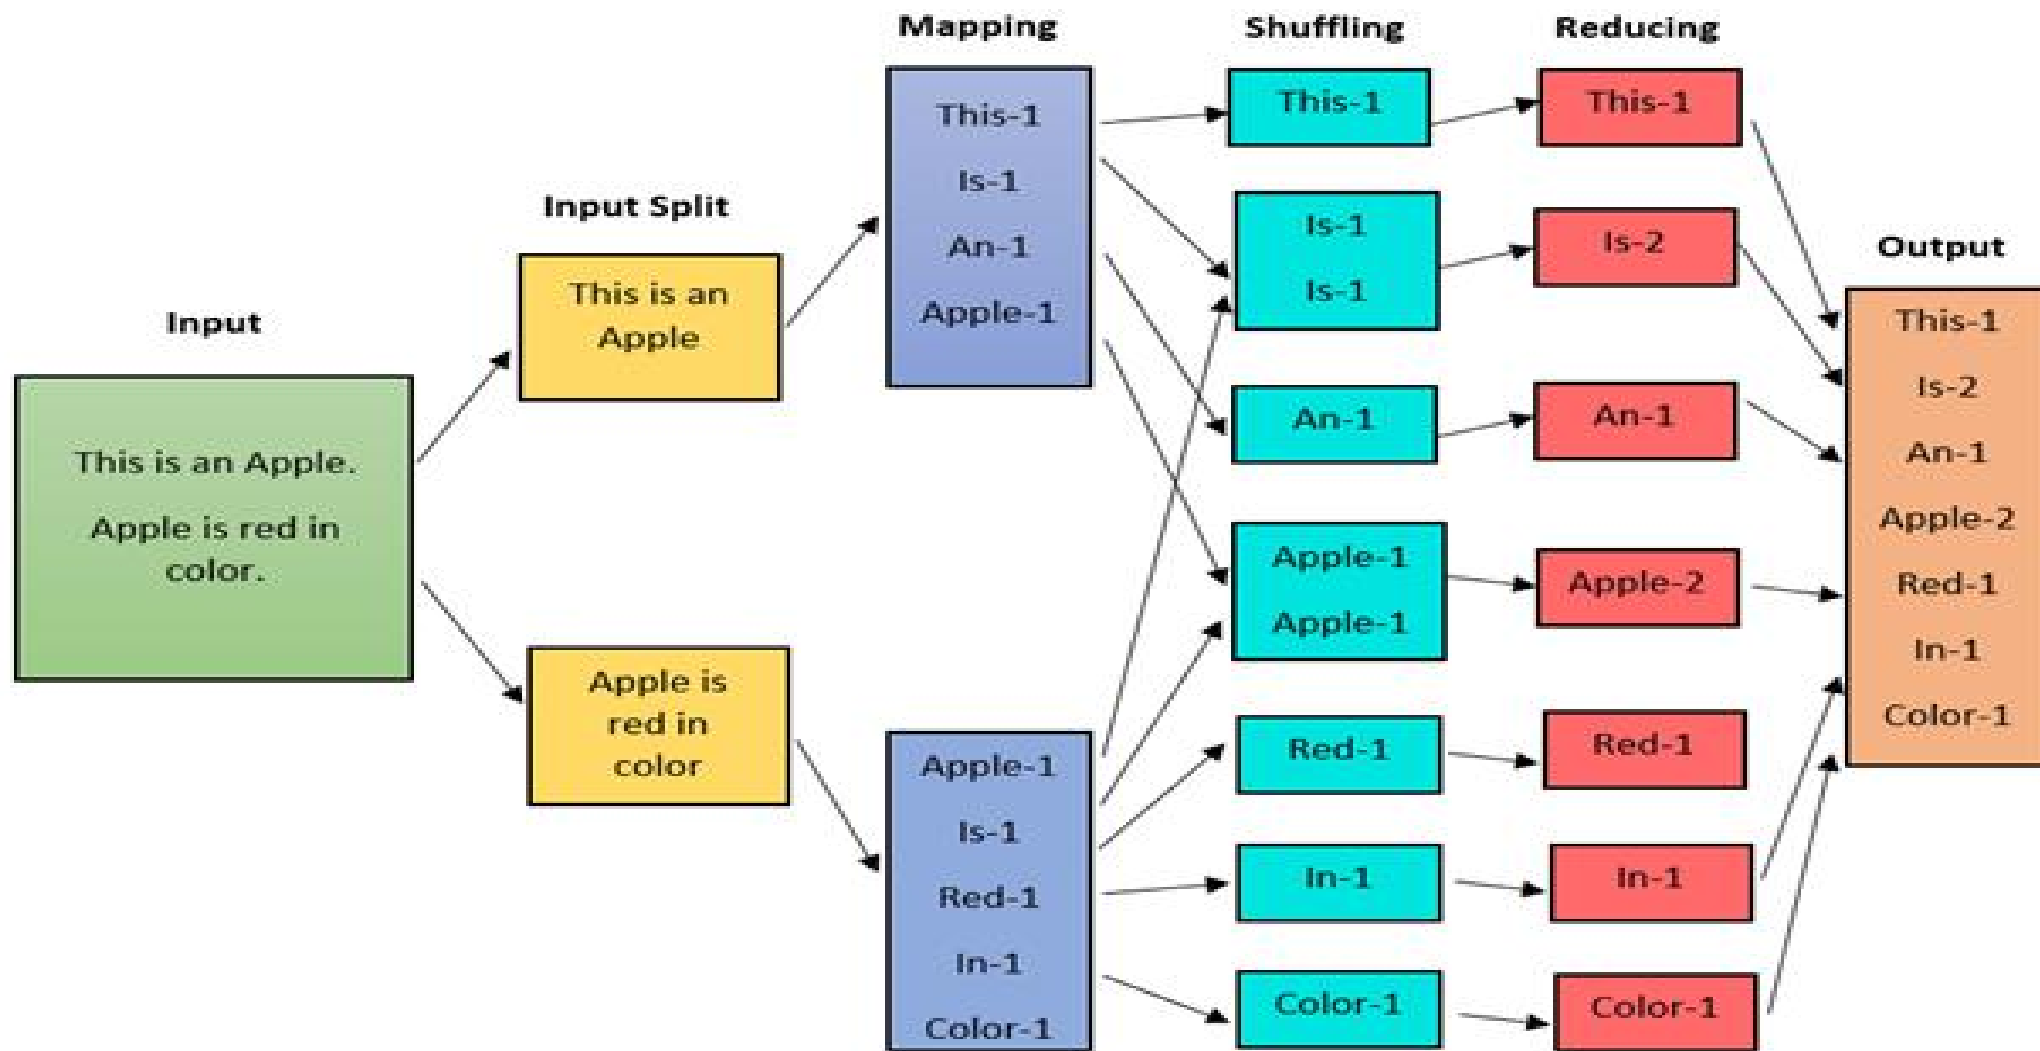
MapReduce- Key Idea

Key idea: Programmers specify two functions:

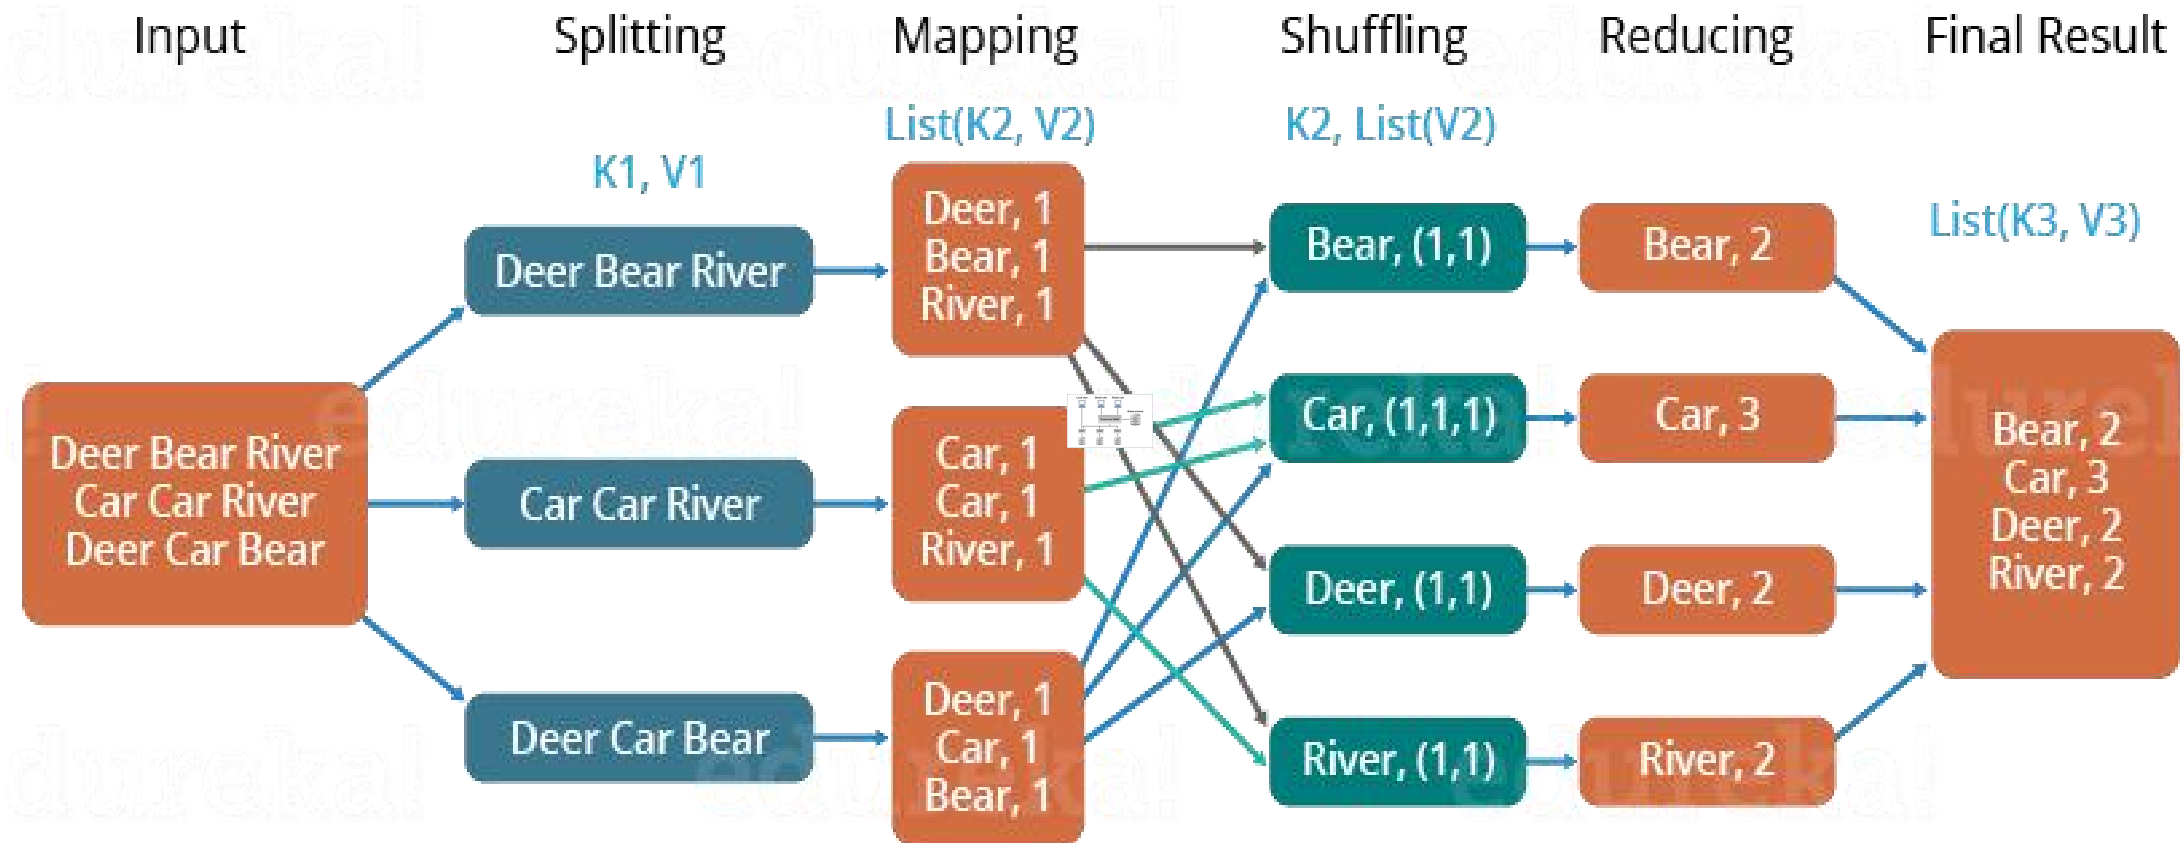
- *map* $(k, v) \rightarrow \langle k', v' \rangle^*$
- *reduce* $(k', v') \rightarrow \langle k', v' \rangle^*$
- *All values with the same key are sent to the same reducer*

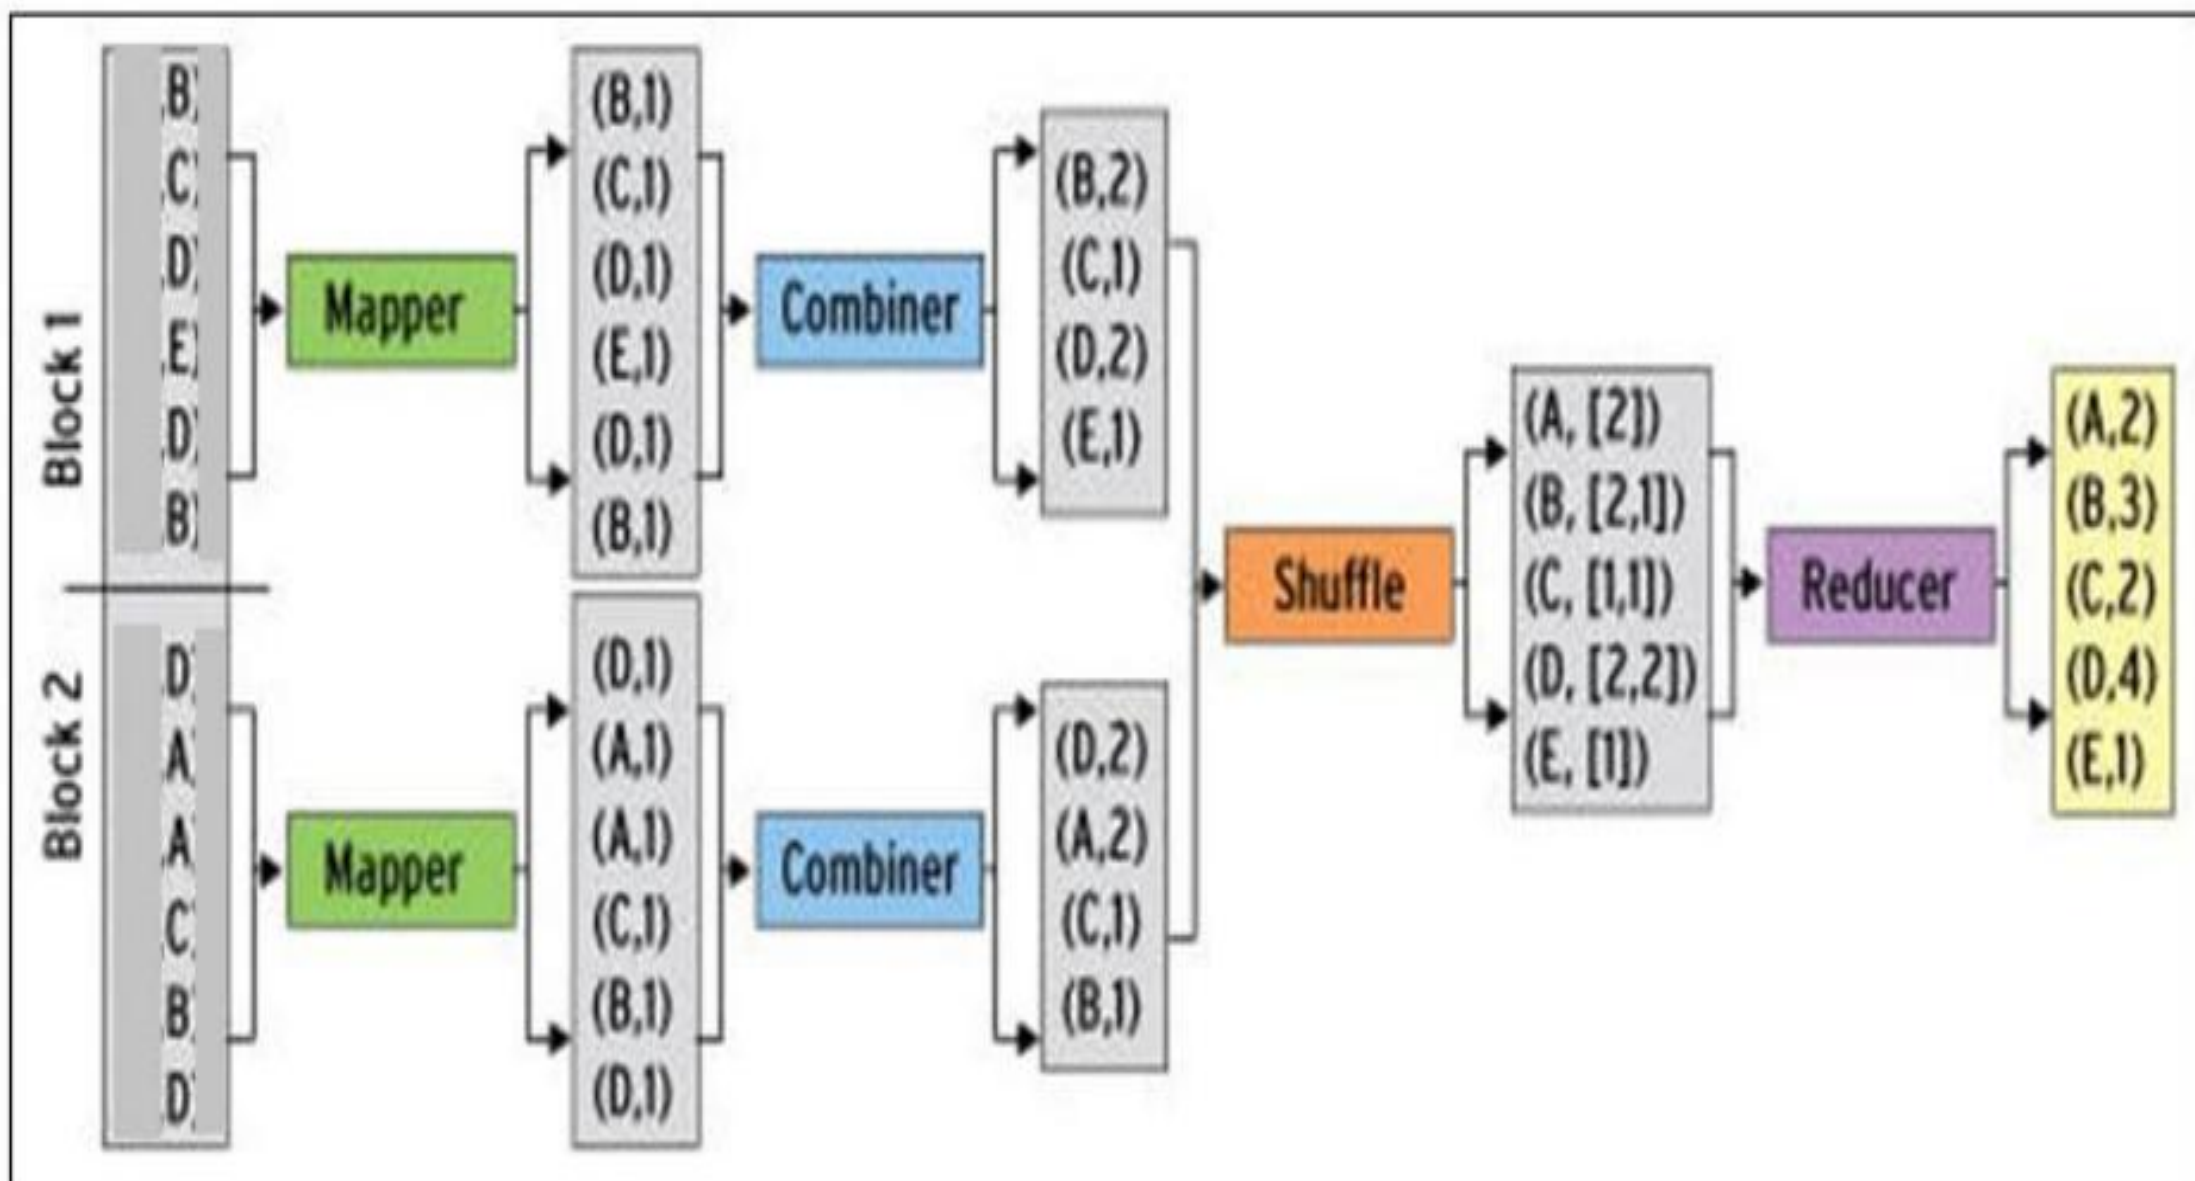
MapReduce – Word Count

- We have a huge text document
- Count the number of times each distinct word appears in the file
- Sample application:
 - Analyze web server logs to find popular URLs



The Overall MapReduce Word Count Process





Map-Reduce Environment is responsible for

Partitioning the input data

Scheduling the program's execution across a set of machines

- *Assigns workers to map and reduce tasks*

Handles synchronization

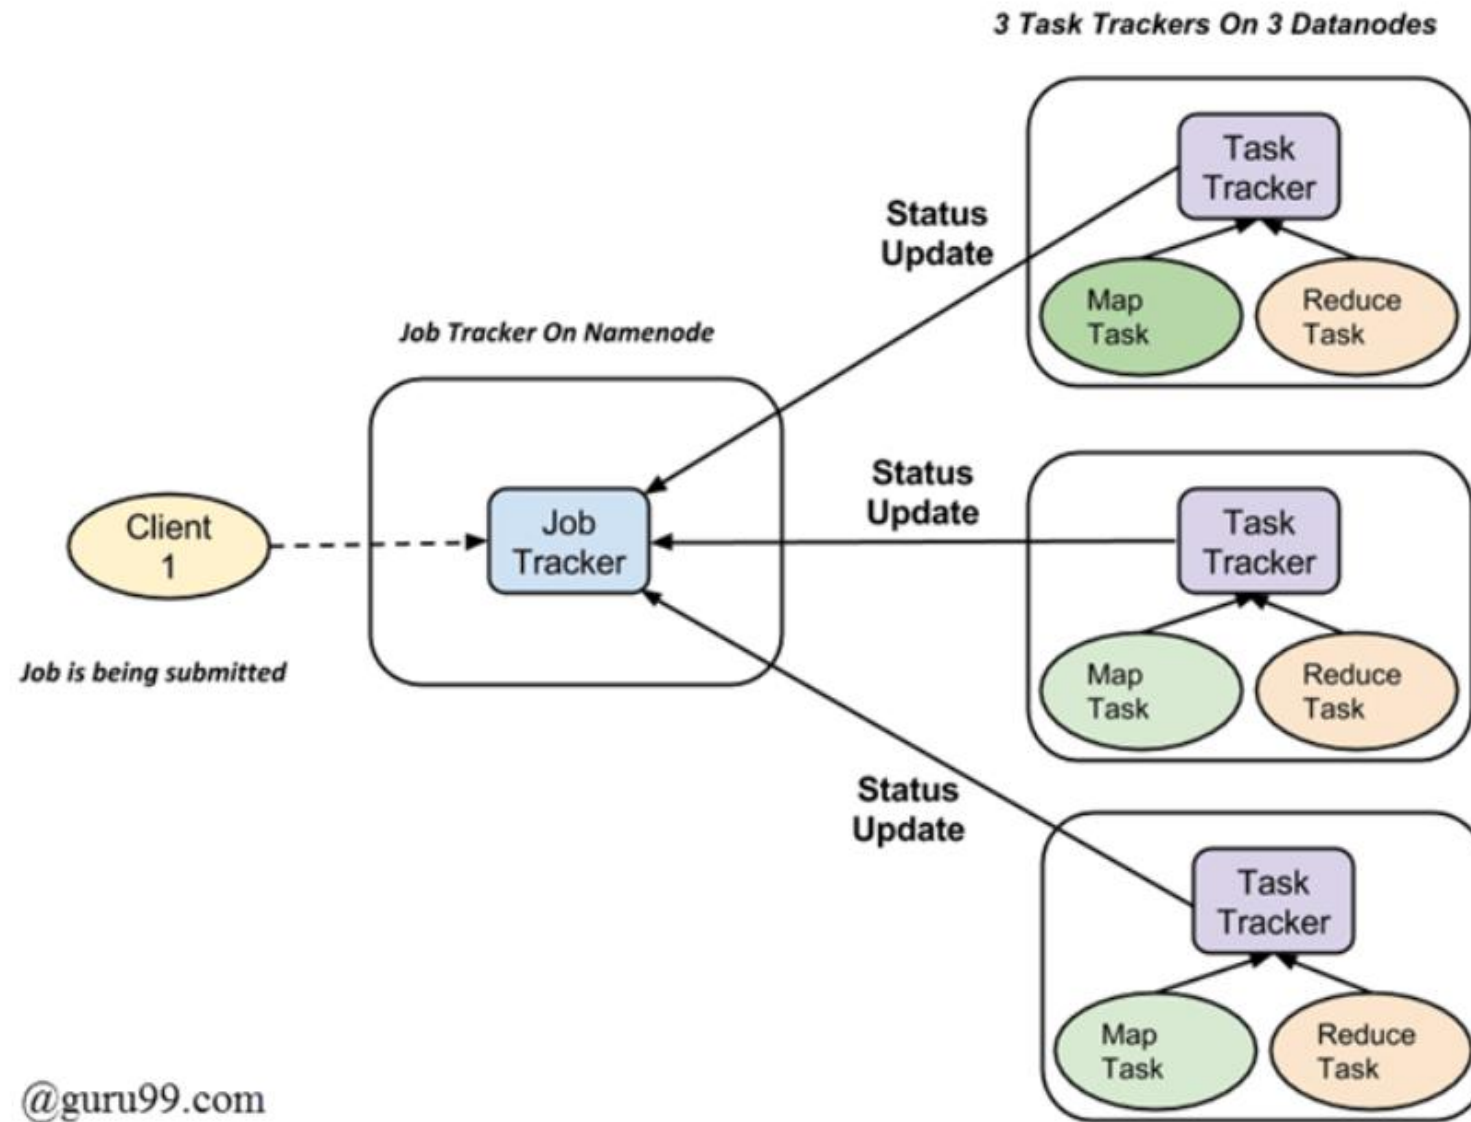
- *Gathers, sorts, and shuffles intermediate data*

Handling machine failures

- *Detects worker failures and restarts*

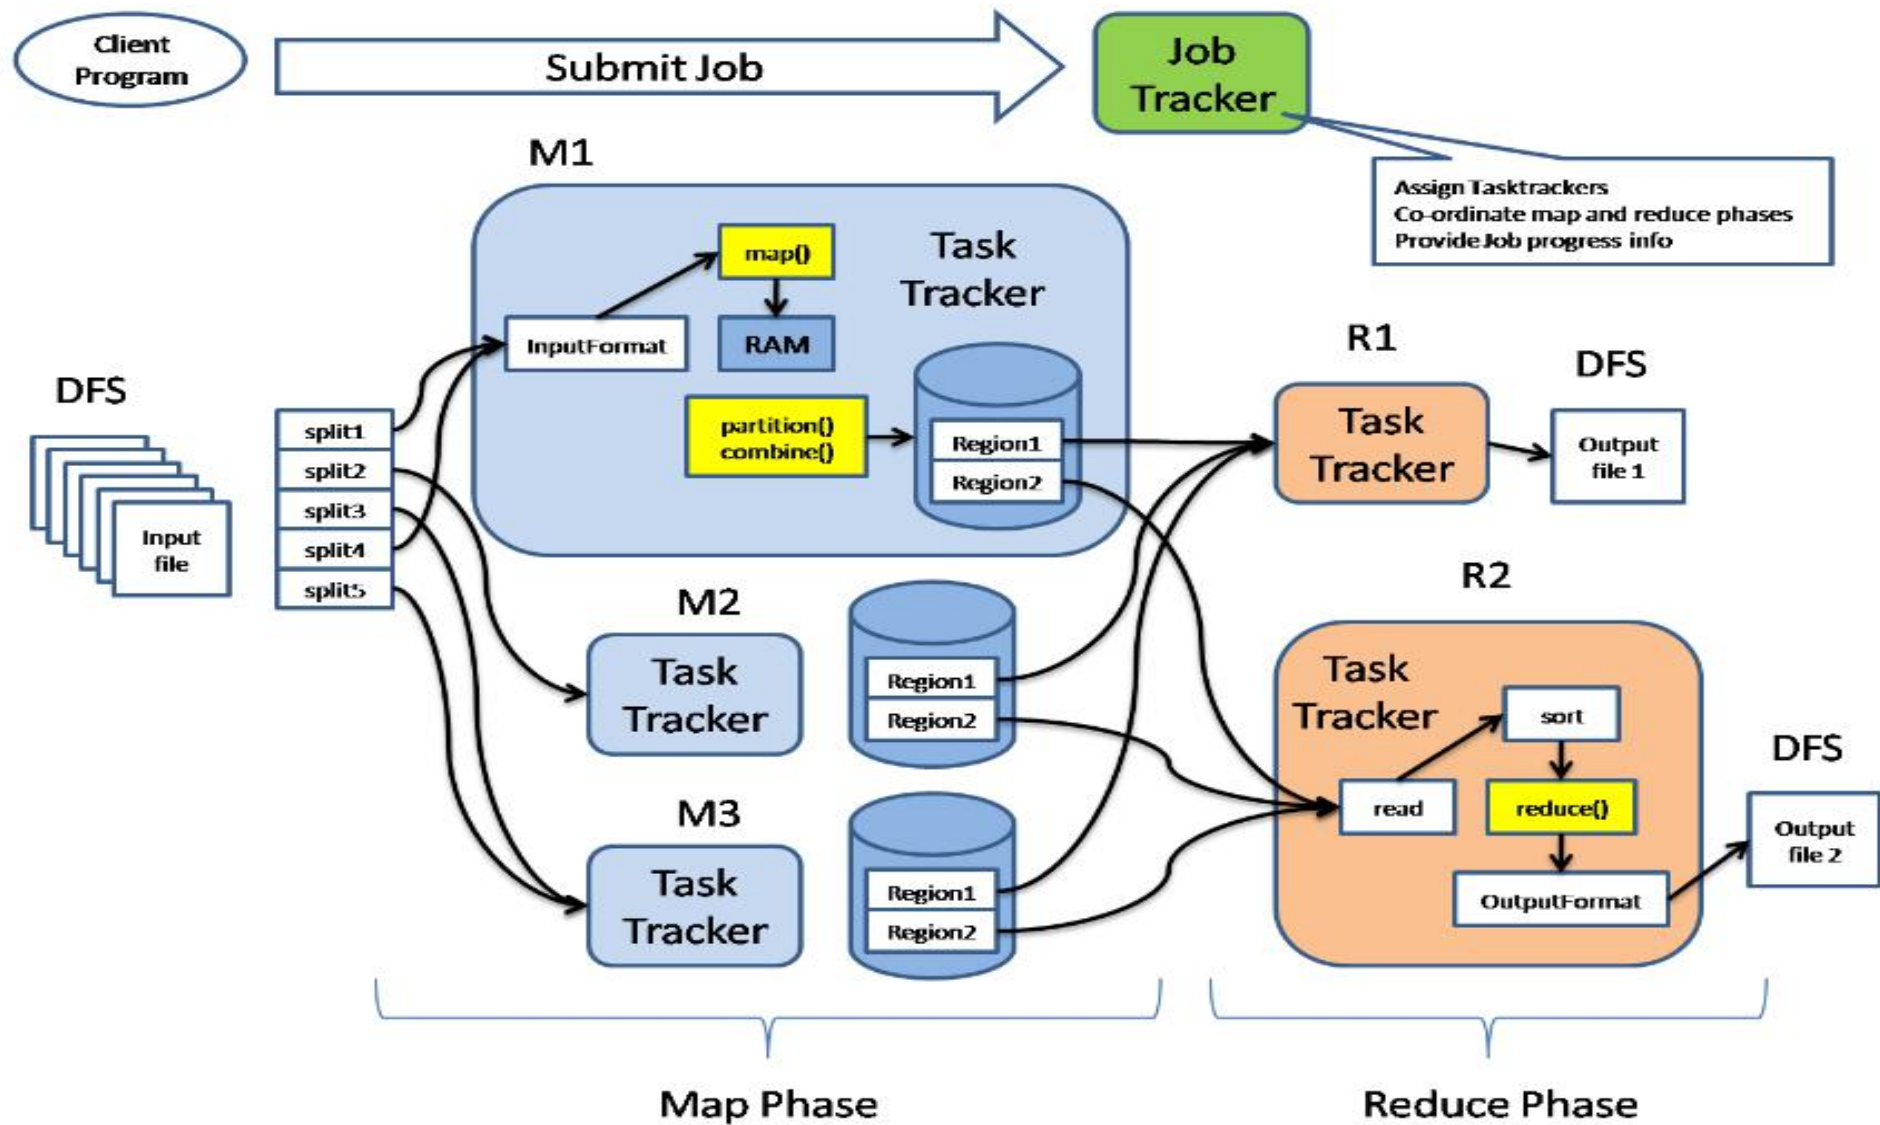
Managing required inter-machine communication

- The complete execution process (Map & Reduce task) is controlled by
 - Jobtracker
 - Task Trackers
- **Job Tracker:** This tracker plays the role of **scheduling jobs** and **tracking all jobs** assigned to the task tracker. Act like a master.
- **Task Tracker:** This tracker plays the role of **tracking tasks** and **reporting the status of tasks** to the job tracker. Act like a slave.
- **Task tracker** periodically sends 'heartbeat' signal to the **Jobtracker** so as to **notify** him of the **current state of the system**.
- In the event **of task failure**, the **job tracker can reschedule** it on a different task tracker.



@guru99.com

How Hadoop MapReduce Works



Mapper.py

```
#!/usr/bin/env python
```

```
import sys
```

```
# Read each line from stdin
```

```
# Generate the count for each word
```

```
for word in words:
```

```
# Write the key-value pair to stdout to be processed by  
# the reducer.
```

```
# The key is anything before the first tab character and the  
#value is anything after the first tab character.
```

```
print '{0}\t{1}'.format(word, 1)
```

Reducer.py

```
#!/usr/bin/env python

import sys

curr_word = None
curr_count = 0

# Process each key-value pair from the mapper
for line in sys.stdin:

    # Get the key and value from the current line
    word, count = line.split('\t')

    # Convert the count to an int
    count = int(count)

    # If the current word is the same as the previous word,
    # increment its count, otherwise print the words count
    # to stdout
    if word == curr_word:
        curr_count += count
    else:

        # Write word and its number of occurrences as a key-value
        # pair to stdout
        if curr_word:
            print '{0}\t{1}'.format(curr_word, curr_count)

        curr_word = word
```

```
curr_count = count
```

```
# Output the count for the last word
```

```
if curr_word == word:
```

```
    print '{0}\t{1}'.format(curr_word, curr_count)
```


Word Count using MapReduce

```
map(key, value):
```

```
// key: document name; value: text of the document
```

```
  for each word w in value:
```

```
    emit(w, 1)
```

```
reduce(key, values):
```

```
// key: a word; value: an iterator over counts
```

```
  result = 0
```

```
  for each count v in values:
```

```
    result += v
```

```
  emit(key, result)
```

```
from mrjob.job import MRJob

class MRWordCount(MRJob):

    def mapper(self, _, line):
        for word in line.split():
            yield(word, 1)

    def reducer(self, word, counts):
        yield(word, sum(counts))

if __name__ == '__main__':
    MRWordCount.run()
```

```
map(key, value):
```

```
// key: document name; value: text of
the document
```

```
for each word w in value:
```

```
    emit(w, 1)
```

```
reduce(key, values):
```

```
// key: a word; value: an array counts
result = 0
```

```
for each count v in values:
```

```
    result += v
```

```
emit(key, result)
```


Example: Inverted Index

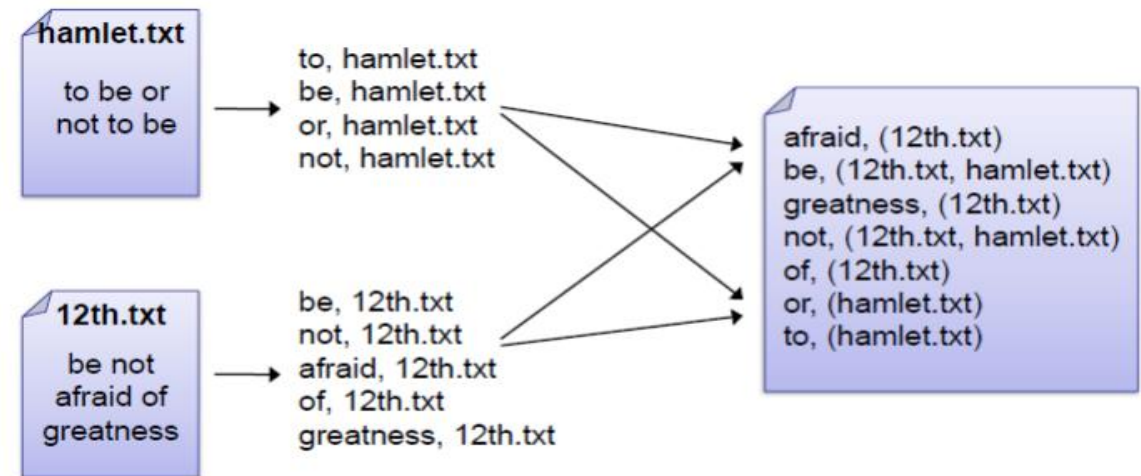
- This was the original Google's usecase
- Generate an inverted index of words from a given set of files

- **Map:**

- parses a document and emits $\langle \text{word}, \text{docId} \rangle$ pairs

- **Reduce:**

- takes all pairs for a given word, sorts the docId values, and emits a $\langle \text{word}, \text{list}(\text{docId}) \rangle$ pair



Example: Language modeling

- Statistical machine translation:
 - *Need to count number of times every 5-word sequence occurs in a large corpus of documents*

Map

- *extract (5-word sequence, count) from document*

Reduce

- *combine counts*

Example: Distributed Grep

- Find all occurrences of the **given pattern** in a very large set of files.

Map:

- *Apply grep on assigned documents*
- *Emit list of documents that contain term*

Reduce:

- *Merge lists*

Dealing with Failures

Map worker failure

- *Map tasks completed or in-progress at worker are restarted*

Reduce worker failure

- *Only in-progress tasks are restarted*

Master failure

- *MapReduce task is aborted and client is notified*

References

- <https://www.guru99.com/introduction-to-mapreduce.html>
- <https://www.section.io/engineering-education/understanding-map-reduce-in-hadoop/>
- <https://www.geeksforgeeks.org/introduction-to-hadoop-distributed-file-systemhdfs/>
- <https://ia600201.us.archive.org/7/items/HadoopInPractice/Hadoop%20in%20Practice.pdf>
- <https://pepa.holla.cz/wp-content/uploads/2016/10/hadoop-with-python.pdf>