

MapReduce

lec 05

K-Mean

- K-mean is **unsupervised machine learning algorithm** which groups the unlabelled dataset into different clusters based on their similarity.
- k-means clustering **divides data into a predefined number of clusters.**
- It is a **centroid-based algorithm**, where **each cluster is associated with a centroid.** The main aim of this algorithm is to **minimize the sum of distances between the data point and their corresponding clusters.**

K-mean Algorithm

- **Initialization:** Choose K initial centroids randomly from the dataset. These centroids are the initial cluster centers.
- **Assignment:** For each data point in the dataset, calculate the distance (usually Euclidean distance) between that point and each of the K centroids. Assign the data point to the cluster whose centroid is closest to it.
- **Update:** Recalculate the centroids of each cluster by taking the mean of all data points assigned to that cluster.

K-mean Algorithm

- **Repeat:** Repeat the assignment and update steps until a convergence criterion is met. Common convergence criteria include when the centroids no longer change significantly, or when a fixed number of iterations is reached.
- **Result:** The final centroids represent the centers of the K clusters, and each data point belongs to the cluster associated with the nearest centroid.

Example

- $dist = \sqrt{(\mu - x)^2}$

Dataset		centroid 1 (3)	centroid 2(11)	cluster
1	2	1	9	1
2	4	1	7	1
3	10	7	1	2
4	12	9	1	2

cluster 1={2,4}

cluster 2={10,12}

New centroid 1: $(2+4)/2=3$

new centroid 2: $(10+12)/2= 11$

the centroid remain same so the results is

cluster 1={2,4,3} with centroid 3

cluster 2={10,12,11} with centroid 11

Parallel K-mean using MapReduce

- **Mapper:**

- assign data points to closest cluster center.

$$z_i \leftarrow \arg \min_j ||\mu_j - \mathbf{x}_i||_2^2$$

Map: For each data point, given $(\{\mu_j\}, \mathbf{x}_i)$, emit(z_i, \mathbf{x}_i)

- **Reducer:**

- revise cluster center as mean of assigned observation.

$$\mu_j = \frac{1}{n_j} \sum_{i: z_i=k} \mathbf{x}_i$$

Reduce: Average over all points in cluster j ($z_i=k$)

Mapper

$$z_i \leftarrow \arg \min_j ||\mu_j - \mathbf{x}_i||_2^2$$

map($[\mu_1, \mu_2, \dots, \mu_k], \mathbf{x}_i$)

$$z_i \leftarrow \arg \min_j ||\mu_j - \mathbf{x}_i||_2^2$$

emit(z_i, \mathbf{x}_i)

Reducer

revise cluster center as mean of assigned observation

$$\mu_j = \frac{1}{n_j} \sum_{i: z_i = k} \mathbf{x}_i$$

```
reduce(j, x_in_cluster j : [x1, x3,..., ])
```

```
    sum = 0
```

```
    count = 0
```

```
    for x in x_in_cluster j
```

```
        sum += x
```

```
        count += 1
```

```
    emit(j, sum/count)
```


K-mean using MapReduce

```
from mrjob.job import MRJob  
import numpy as np
```

```
class ClusterAssignmentMapReduce(MRJob):
```

```
    def configure_args(self):
```

```
        super(ClusterAssignmentMapReduce, self).configure_args()
```

```
        self.add_file_arg('--clusters', help='File containing cluster centroids')
```

```
def mapper_init(self):
    # Load cluster centroids from the file
    with open(self.options.clusters, 'r') as f:
        self.cluster_centroids = [float(line.strip()) for line in f]

def mapper(self, _, line):
    # Parse the input data point
    data_point = float(line.strip())

    # Find the nearest cluster for the data point
    nearest_cluster = min(self.cluster_centroids, key=lambda centroid: abs(centroid -
data_point))

    # Emit the nearest cluster and the data point
    yield nearest_cluster, data_point
```

```
def reducer(self, cluster, data_points):  
    # Collect and emit data points for each cluster  
    yield cluster, list(data_points)
```

```
if __name__ == '__main__':  
    ClusterAssignmentMapReduce.run()
```